



Production ready Data Science



About this Talk

- Interactively build a model that classifies internet posts
- Get a glimpse of production-readiness
- Learn about Luigi pipelines and their main components
- Write your production ready pipeline
- Get an overview of luigis modules

Building the model interactively

Download dataset

```
In [11]: !mkdir -p dataset
!curl -L -o dataset/dataset.zip http://plainpixels.work/resources/datasets/reddit_
ds_got.zip
!unzip -u -q -d dataset/reddit dataset/dataset.zip
```

% Total		% Received		% Xferd		Average Speed		Time	Time	Time	Curren
											t
						Dload	Upload	Total	Spent	Left	Speed
100	318k	100	318k	0	0	217k	0	0:00:01	0:00:01	--:--:--	217k

```
In [12]: !ls dataset/reddit
```

```
02-19-2018 02-20-2018 02-21-2018 02-22-2018 __MACOSX training
```

Explore the data

```
In [38]: dataset = pandas.read_csv("dataset/reddit/training/data.csv",
                                   encoding='utf-8',
                                   sep=';').fillna('')

print(dataset.count())
print(dataset.describe())
dataset.head()
```

```
title          1000
selftext       1000
subreddit      1000
dtype: int64

              title  selftext  subreddit
count              1000      1000        1000
unique              999        708           2
top  Hiring Data Scientists      datascience
freq              2        293          500
```

Out [38]:

	title	selftext	subreddit
0	Weekly 'Entering & Transitioning' Thread. Ques...	**Welcome to the very first 'Entering & Transi...	datascience
1	Data Science in Fashion		datascience
2	Evaluating our Startup using 3 versions of the...		datascience
3	Best place to read other people's code?	Hello,\n\nI'm starting to learn data science a...	datascience
4	What are the best practices for downloading da...	Downloading data from an API is a really commo...	datascience


```
In [62]: ds = " ".join(selftext for selftext
                        in dataset[dataset["subreddit"]=="datascience"].selftext)
ds_wordcloud = WordCloud(min_font_size=8,
                          max_words=100,
                          background_color="white").generate(ds)

plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(ds_wordcloud, interpolation='bilinear')
plt.axis("off")
plt.tight_layout(pad = 0)

plt.show()
```



In [59]:

```
tokenizer = nltk.tokenize.RegexpTokenizer(r'\w+')
stopwords = nltk.corpus.stopwords.words('english')
stemmer = nltk.SnowballStemmer("english")

def clean_words(post):
    tokenized = tokenizer.tokenize(post["title"] + " " + post["selftext"])
    lowercase = [word.lower() for word in tokenized]
    filtered = [word for word in lowercase if word not in stopwords]
    stemmed = [stemmer.stem(word) for word in filtered]
    return " ".join(stemmed)

dataset["cleaned_words"] = dataset.apply(clean_words, axis=1)

dataset.describe()
```

Out [59]:

	title	selftext	subreddit	cleaned_words
count	1000	1000	1000	1000
unique	999	708	2	1000
top	Hiring Data Scientists		datascience	month month member select hello frand got inte...
freq	2	293	500	1


```
In [64]: ds_clean = " ".join(selftext for selftext
                                in dataset[dataset["subreddit"]=="datascience"].cleaned_words)
ds_wordcloud_clean = WordCloud(min_font_size=8,
                                max_words=100,
                                background_color="white").generate(ds_clean)

plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(ds_wordcloud_clean, interpolation='bilinear')
plt.axis("off")
plt.tight_layout(pad = 0)

plt.show()
```



```
In [65]: with open("dataset/reddit/training/clean.csv", "w") as out:
          dataset[["cleaned_words", "subreddit"]].to_csv(out, encoding='utf-8', ind
ex=False, sep=';')
```

Build the model

```
In [ ]: # Lade die bereinigten Daten
df = sql.read.format("com.databricks.spark.csv") \
    .option("header", "true") \
    .option("delimiter", ";") \
    .load("dataset/reddit/training/clean.csv")

# Den Klassifikator trainieren
labeled = df.withColumn("label", df.subreddit.like("datascience").cast("double"))
train_set, test_set = labeled.randomSplit([0.8, 0.2])
tokenizer = Tokenizer().setInputCol("cleaned_words").setOutputCol("tokenized")
hashing = HashingTF().setNumFeatures(1000).setInputCol("tokenized").setOutputCol(
    "features")
decision_tree = DecisionTreeClassifier()
pipeline = Pipeline(stages=[tokenizer, hashing, decision_tree])
model = pipeline.fit(train_set)
```

It works, now DEPLOY it!

Integration
Testing

Validate
Quality

Reproducible
Training

Reviews

Test
Regressions

CI/CD

Scheduling

Monitoring

Alerting

Proper
Logging

Collect
Metrics

Unit Testing



A bit about Luigi

Luigi helps to stitch long running tasks together into pipelines

It contains a wide toolbox of task templates (e.g. Hive, Pig, Spark, Python)

How to compose workflows?

A workflow consists of Tasks, Targets and Parameters

Targets correspond to a file or a database entry or some other kind of checkpoint

Tasks consume Targets of other tasks, run a computation, then output a target

Parameters take care of task parameterization

How would our Workflow look like?

- Download the dataset
- Clean the data
- Check for existing classifier
- Classify the posts and save result

